

CHAPTER 5

Buffer Overflows

Slides adapted from "Foundations of Security: What Every Programmer Needs To Know" by Neil Daswani, Christoph Kern, and Anita Kesavan (ISBN 1590597842; <http://www.foundationsofsecurity.com>). Except as otherwise noted, the content of this presentation is licensed under the Creative Commons 3.0 License.





Agenda

- Buffer overflows: attacker hijacks machine
 - Attacker injects malicious code into program
- Preventable, but common
(50% CERT advisories
decade after Morris Worm)
- Fixes: Safe string libraries,
StackGuard, Static Analysis
- Other Types of Overflows: Heap, Integer, ...

6.1. Anatomy of a Buffer Overflow

- *Buffer*: memory used to store user input, has fixed maximum size
- *Buffer overflow*: when user input exceeds max buffer size
 - Extra input goes into unexpected memory locations



6.1.1. A Small Example

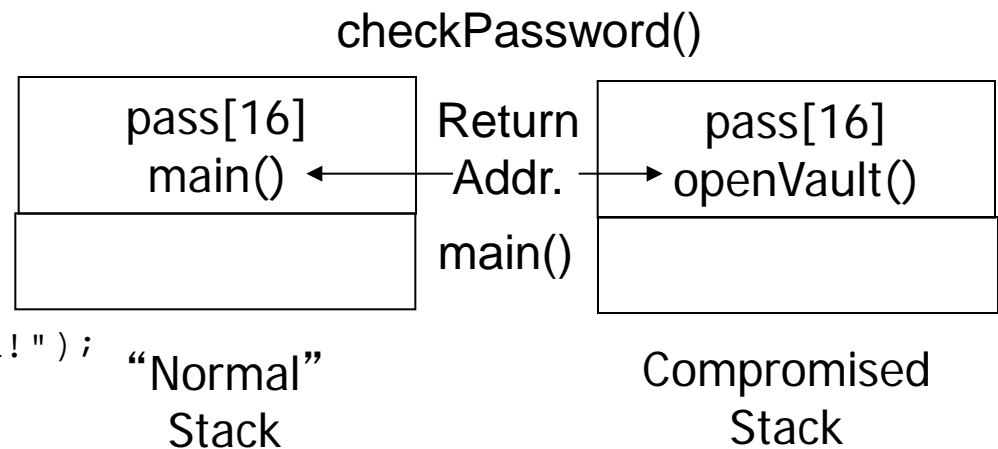

- Malicious user enters > 1024 chars, but buf can only store 1024 chars; extra chars overflow buffer

```
void get_input() {  
    char buf[1024];  
    gets(buf);  
}  
void main(int argc, char*argv[]){  
    get_input();  
}
```



6.1.2. A More Detailed Example

```
1 int checkPassword() {
2     char pass[16];
3     bzero(pass, 16); // Initialize
4     printf ("Enter password: ");
5     gets(pass);
6     if (strcmp(pass, "opensesame") == 0)
7         return 1;
8     else
9         return 0;
10 }
11
12 void openVault() {
13     // Opens the vault
14 }
15
16 main() {
17     if (checkPassword()) {
18         openVault();
19         printf ("Vault opened!");
20     }
21 }
```



6.1.2. checkPassword () Bugs

- *Execution stack*: maintains current function state and address of return function
- *Stack frame*: holds vars and data for function
- Extra user input (> 16 chars) overwrites return address
 - *Attack string*: 17-20th chars can specify address of `openVault ()` to bypass check
 - Address can be found with source code or binary

6.1.2. Non-Executable Stacks Don't Solve It All

- Attack could overwrite return address to point to newly injected code
- NX stacks can prevent this, but not the vault example (jumping to an existing function)
- *Return-into-libc attack*: jump to library functions
 - e.g. `/bin/sh` or `cmd.exe` to gain access to a command shell (*shellcode*) and complete control

6.1.3. The `safe_gets()` Function

```
#define EOLN '\n'
void safe_gets (char *input, int max_chars) {
    if ((input == NULL) || (max_chars < 1)) return;
    if (max_chars == 1) { input[0] = 0; return; }
    int count = 0;
    char next_char;
    do {
        next_char = getchar(); // one character at a time
        if (next_char != EOLN)
            input[count++] = next_char;
    } while ((count < max_chars-1) && // leave space for null
              (next_char != EOLN));
    input[count]=0;
}
```

- Unlike `gets()`, takes parameter specifying max chars to insert in buffer
- Use in `checkPassword()` instead of `gets()` to eliminate buffer overflow vulnerability

```
5 safe_gets(pass, 16);
```


6.2. Safe String Libraries

- C – Avoid (no bounds checks): `strcpy()`, `strcat()`, `sprintf()`, `scanf()`
- Use safer versions (with bounds checking): `strncpy()`, `strncat()`, `fgets()`
- Microsoft's *StrSafe*, Messier and Viega's *SafeStr* do bounds checks, null termination
- Must pass the right buffer size to functions!
- C++: STL string class handles allocation
- Unlike compiled languages (C/C++), interpreted ones (Java/C#) enforce type safety, raise exceptions for buffer overflow

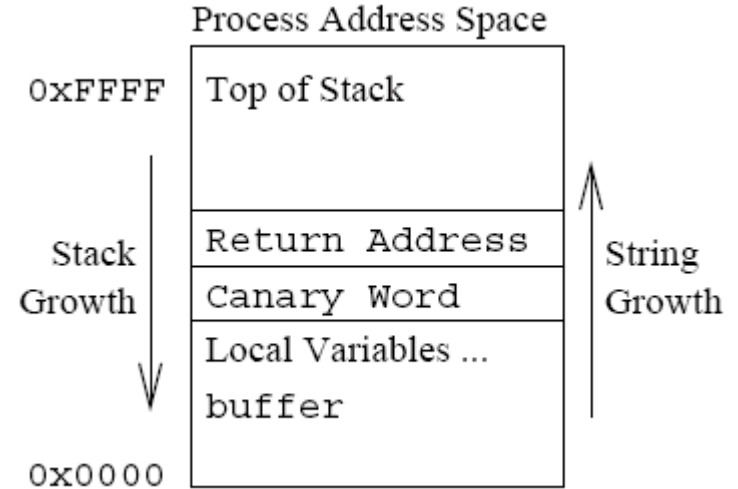


6.3. Additional Approaches

- Rewriting old string manipulation code is expensive, any other solutions?
- StackGuard/canaries (Crispin Cowan)
- Static checking (e.g. Coverity)
- Non-executable stacks
- Interpreted languages (e.g., Java, C#)

6.3.1. StackGuard

- *Canary*: random value, unpredictable to attacker
- Compiler technique: inserts *canary* before return address on stack
- Corrupt Canary: code halts program to thwart a possible attack



Source: C. Cowan et. al., StackGuard,



6.3.2. Static Analysis Tools

- *Static Analysis*: analyzing programs without running them
- Meta-level compilation
 - Find security, synchronization, and memory bugs
 - Detect frequent code patterns/idioms and flag code anomalies that don't fit
- Ex: Coverity, Fortify, Ounce Labs, Klockwork
 - Coverity found bugs in Linux device drivers



6.4. Performance

- Mitigating buffer overflow attacks incurs little performance cost
- Safe str functions take slightly longer to execute
- StackGuard canary adds small overhead
- But performance hit is negligible while security payoff is immense

6.5. Heap-Based Overflows

- Ex: `malloc()` in C provides a fix chunk of memory on the heap
- Unless `realloc()` called, attacker could also overflow heap buffer (fixed size), overwrite adjacent data to modify control path of program
- Same fixes: bounds-checking on input



6.6. Other Memory Corruption Vulnerabilities

- *Memory corruption vulnerability*: Attacker exploits programmer memory management error
- Other Examples
 - Format String Vulnerabilities
 - Integer Overflows
 - Used to launch many attacks including buffer overflow
 - Can crash program, take full control

6.6.1. Format String Vulnerabilities

- Format String in C directs how text is formatted for output: e.g. %d, %s
 - Can contain info on # chars (e.g. %10s)

```
void format_warning (char *buffer,  
                    char *username, char *message) {  
    sprintf (buffer, "Warning: %10s -- %8s",  
            message, username);  
}
```

- If message or username greater than 10 or 8 chars, buffer overflows
 - attacker can input a username string to insert shellcode or desired return address

6.6.2. Integer Overflows (1)

- Exploits range of value integers can store
 - Ex: signed four-byte int stores between -2^{32} and $2^{32}-1$
 - Cause unexpected wrap-around scenarios
- Attacker passes `int` greater than max (positive)
 - > value wraps around to the min (negative!)
 - Can cause unexpected program behavior, possible buffer overflow exploits

6.6.2. Integer Overflows (2)

```
/* Writes str to buffer with offset
characters of blank spaces
preceding str. */
3 void formatStr(char *buffer, int buflen,
4               int offset, char *str, int slen) {
5     char message[slen+offset];
6     int i;
7
8     /* Write blank spaces */
9     for (i = 0; i < offset; i++)
10         message[i] = ' ';
11
12     strncpy(message+offset, str, slen);
13     // offset = 232!?
14     strncpy(buffer, message, buflen);
15     message[buflen-1] = 0;
16     /*Null terminate*/
17 }
```

- Attacker sets
offset = 2^{32} ,
wraps around to
negative values!
 - write outside
bounds of
message
 - write arbitrary
addresses
on heap!



Summary

- Buffer overflows most common security threat!
 - Used in many worms such as Morris Worm
 - Affects both stacks and heaps
- Attacker can run desired code, hijack program
- Prevent by bounds-checking all buffers
 - And/or use StackGuard, Static Analysis...
- Type of Memory Corruption:
 - Format String Vulnerabilities, Integer Overflow, etc...